eleven - Hackathon

Deep Learning Basics

To the attention of the Ingénierie Mathématique & Informatique students
September 6th, 2022







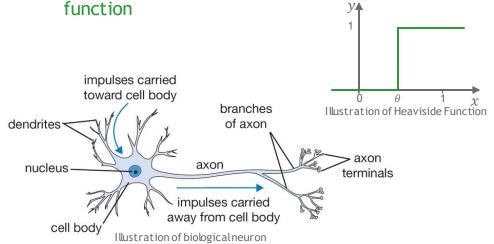
Artificial neural networks are based on biological neurons because the human brain processes information efficiently



Biological neurons

- Basis constituent of the nervous system
- Information is propagated by electric signals through dendrites and axons
- A neuron is activated when the input signal is higher than a threshold

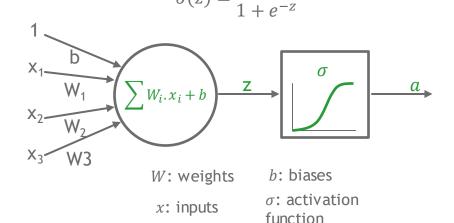
• A neuron is equivalent to a Heaviside step





Artificial neuron

- A simplified model of a biological neuron
- An artificial neuron is a linear function, on which an activation function is added
- The step function is not adapted due to its sensitivity to noise
- The common activation function used is the sigmoid function:

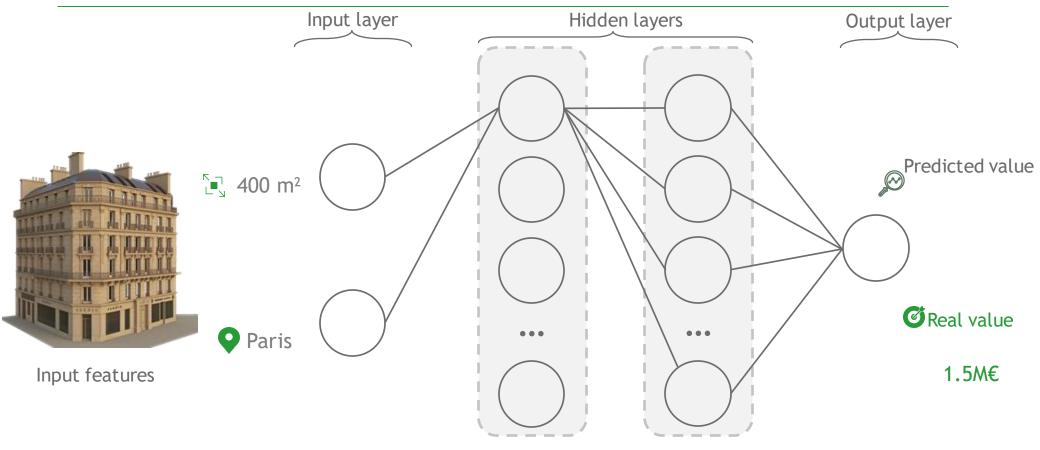




- Each neuron processes a bit of information and passes it to its children
- Overall, the network processes raw information into general concepts

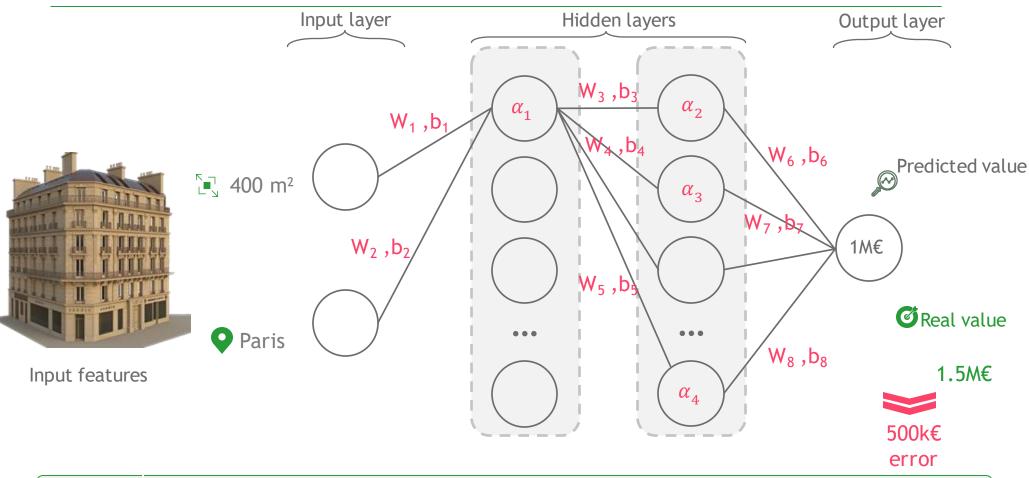
A neural network learn how to imitate a function from a training dataset

Computation process: real estate price prediction example



A neural network learn how to imitate a function from a training dataset

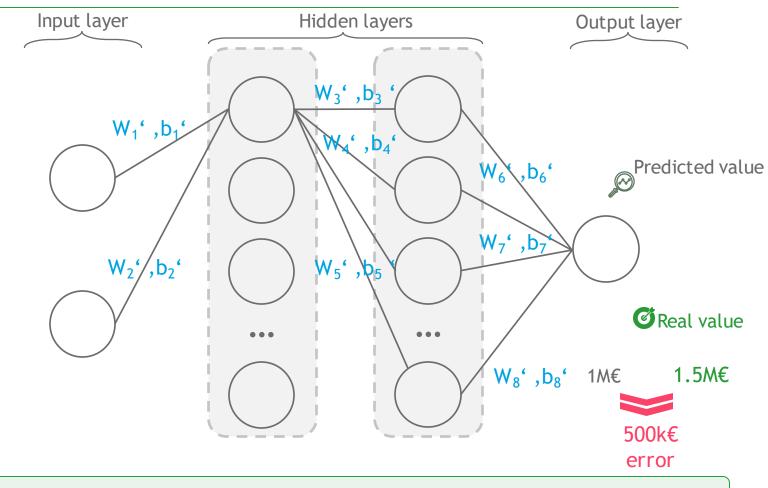
© Computation process: real estate price prediction example



- Note
- This process is called the forward propagation and allows to compute the loss function of the model
- How can we find the model's optimal parameters for this task?

A neural network learn how to imitate a function from a training dataset

Computation process: real estate price prediction example





- This process is called the backward propagation and allows to update the neural network parameters
- Forward and backward propagation are done iteratively until having a high-performance model, by doing a gradient descent to minimize loss

The gradient descent can be controlled by a hyperparameter of the network: the learning rate



Overview of the learning rate

- The learning rate (α) controls the speed at which the model learn
- · It controls the update of the weights
- If the learning rate is high, the weights will be strongly modified at each epoch
- A technique to speed up the gradient descent is the learning rate decay, which consists to slowly reduce the learning rate over time: allows large weight changes at the start of the learning process, and more fine-tuning towards the end



Illustration of the learning rate

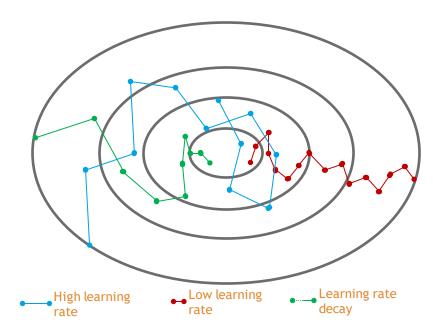


Illustration of the impact of the learning rate on the gradient descent



- The learning rate is equivalent to the length of the step, and it may be the most important hyperparameter to configure our model, and is generally in $[10^{-6}, 1]$
- Adaptive Learning Rates is a learning rate schedule adapting the LR to the performance of the model. Typically, once the performance of the model plateaus, LR can be decreased by a factor of 2

On Python, using an open-source machine learning framework like PyTorch can accelerate the path from research prototyping to production deployment



Deep Learning computation



- PyTorch is defined as an open-source machine learning library for Python
 It is designed for Deep Learning and used for applications such as NLP, Computer Vison



Main Features

- Easy Interface:
 - PyTorch offers easy to use API; hence it is considered to be very simple to operate and runs on Python
- Python usage:
 - · Pytorch can leverage all the services and functionalities offered by the Python environment
- Computational graphs:
 - PyTorch provides a platform which offers dynamic computational graphs. Thus, a user can change them during runtime



Key advantages

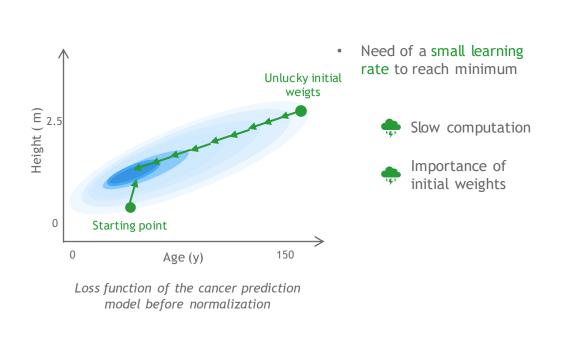
- Easy to debug
- It includes many layers as Torch
- It includes lot of loss functions
- It can be considered as NumPy extension to **GPUs**
- It allows building networks whose structure is dependent on computation itself

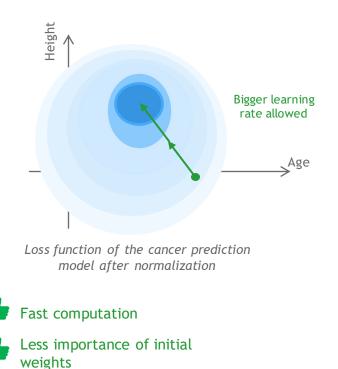
Normalisation throughout the neural network helps ease the optimization process



Why should layers be normalized?

Use case of cancer prediction based on Height and Age







- Use BatchNorm1d to normalize a layer with PyTorch
- Weight initialization can be "random" by default. Otherwise, it can depend on the type of activation function used: for "Sigmoid", a xavier initialization method is usually used

How to prevent Neural Network from overfitting?

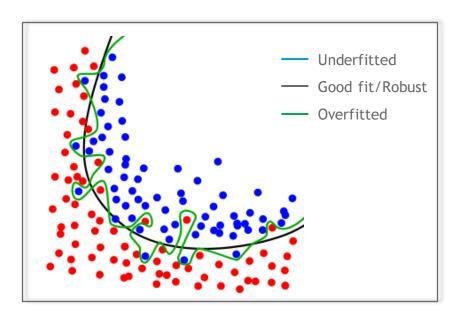
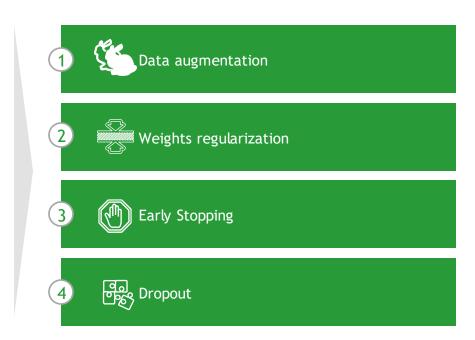
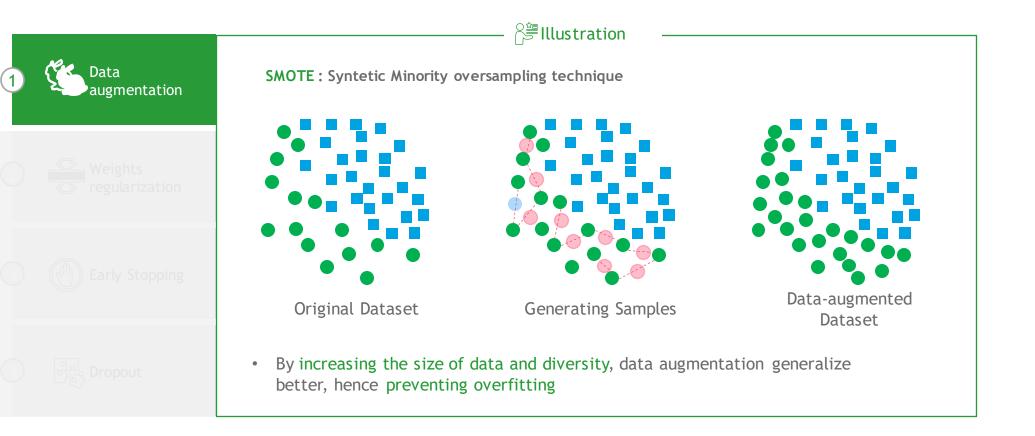


Illustration of a regression problem and overfitting possibilities





How to prevent Neural Network from overfitting?



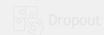


How to prevent Neural Network from overfitting?









Illustration

Lasso regularization

New objective function:

$$\frac{1}{N} \sum_{i=1}^{N} (Y_i - \hat{Y}_i)^2 + \lambda_1 \sum_{j=0}^{n} |\beta_j|$$

2 Ridge regularization

New objective function:

$$\frac{1}{N} \sum_{i=1}^{N} (Y_i - \hat{Y}_i)^2 + \lambda_2 \sum_{j=0}^{n} \beta_j^2$$

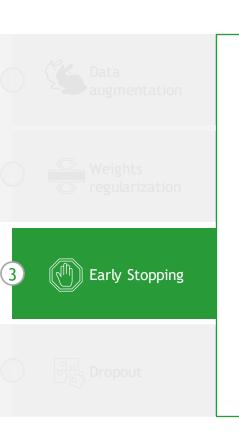
ElasticNet data regularization is the combination of both Lasso and Ridge regularization



torch.optim.SGD(net.parameters(), Lr=0.2, weight_decay=1e-2)



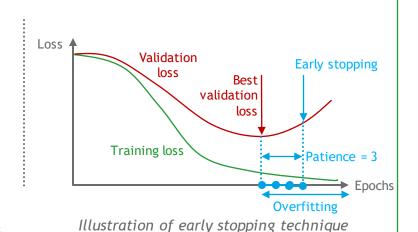
How to prevent Neural Network from overfitting?



Early stopping stops the training when the validation loss is increasing

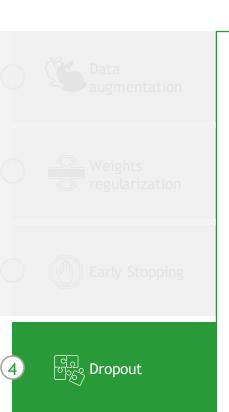
Illustration

- If the validation loss is higher than the current lowest validation lost for a certain number of epochs the training is stopped
- The number of epochs we wait is known as the patience





How to prevent Neural Network from overfitting?



#Illustration

- Not training certain weights at a given update step, leading to greater generalization over the whole network
- The classic formulation of this is dropout where neuron activations are "dropped out" (set to zero) with probability p

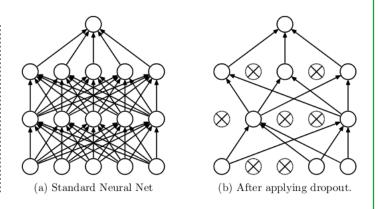


Illustration of dropout technique



torch.nn.Dropout(p=0.5)

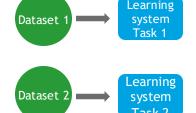
Transfer learning uses knowledge acquired for one task to solve related ones and enables to reduce the data required for the related task



General Overview

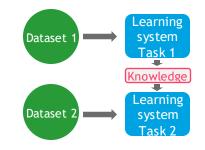
Isolated, single task learning :

- Knowledge is not retained or accumulated
- Learning is performed without considering past learned knowledge in other tasks



Learning of a new task relies on the previous learned tasks:

 Learning process can be faster, more accurate and / or need less training data



There are two variations of transfer learning:

- Same domain but different tasks
- Same task, but different domains



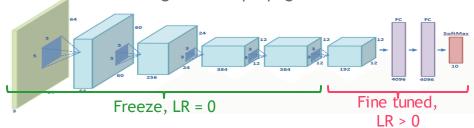
Transfer learning in Deep Learning

Use pre-trained model and freeze most layers of the neural network

 Freeze layer imply that weights are not updated during the backpropagation, the learning rate (LR) is equal to 0

Apply fine tuning to lasts layers in order to adapt the model to the specific task

The weights of fine-tuned are updated during the backpropagation



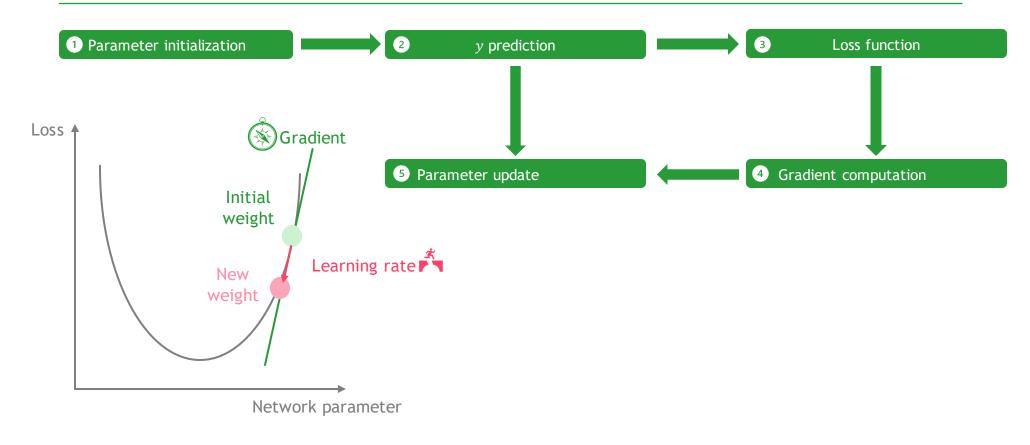
- Keep simple and general features from shallow layers
- · Retrain complex features from deep layers



- Each of the PyTorch domain librarie (torchvision, torchtext) comes with pretrained models (<u>torchvision.models</u>)
- HuggingFaces has a series of pretrained models on many different domains (vision, text, audio) and plenty of different datasets (https://huggingface.co/models)

To conclude, Neural networks are based on human brain to create a powerful model that can learn complex functions

Reminder of what we have learned



Some examples of interesting resources to go further

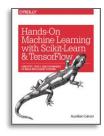
In 1h: Machine Learning for everyone https://vas3k.com/blog/machine_learning/



In 4h: Machine learning Toolbox ("learning by doing")
 https://www.datacamp.com/courses/machine-learning-toolbox



 In 9 Weeks: Hands-On ML with Scikit-Learn & TensorFlow https://www.oreilly.com/library/view/hands-on-machine-learning/9781491962282/



In 11 Weeks: Machine Learning (by Andrew Ng)
 https://www.coursera.org/learn/machine-learning/home/welcome



- And above all
- https://www.coursera.org/specializations/deep-learning

More than 30 Weeks: Deep Learning (Goodfellow et al.)
 http://www.deeplearningbook.org/

