

Training : Computer Vision

Hackathon Ecole des Ponts

06/09/2022



eleven
strategy • data • digital

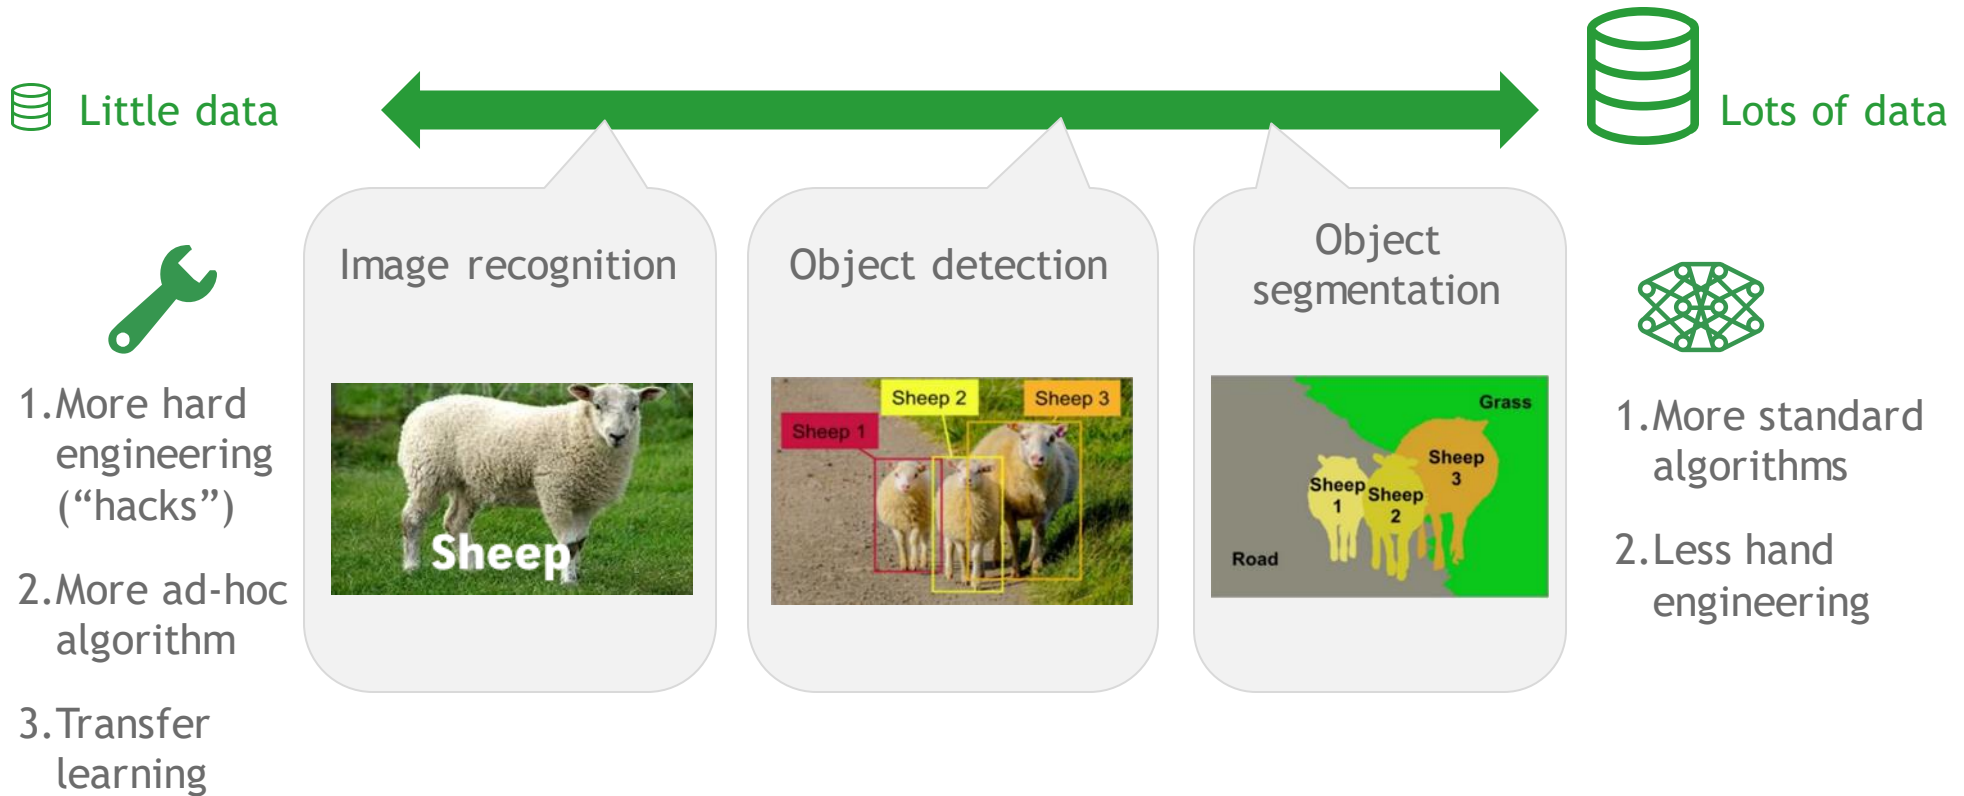


AGENDA



1. **Computer vision approaches**
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm
5. Resources to help you start

The **amount of data available** is a critical element to consider in an image-based application



Sources of knowledge

1. Labeled data (e.g. image → dog & cat)
2. Hand engineered features, network architecture, etc.

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

Image Pre-processing

Image scaling

Contrast
enhancement

Cropping and
padding

Noise reduction




(a)



(b)



(c)



(d)

(a)Original image

(b)Resized

(c)Grayed

(d)Blured

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)




Image Pre-processing

Image scaling
Contrast
enhancement
Cropping and
padding
Noise reduction



Feature extraction

Lines/Edge detection
Corners/Blobs/Points
(SIFT/SURF, PCA,
Watershed)



Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

Image Pre-processing

Image scaling
Contrast
enhancement
Cropping and
padding
Noise reduction



Feature extraction

Lines/Edge detection
Corners/Blobs/Points
(SIFT/SURF, PCA,
Watershed)



Classification algorithm

CNN



Cat

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

Image Pre-processing

Image scaling
Contrast enhancement
Cropping and padding
Noise reduction



Feature extraction


Lines/Edge detection
Corners/B
(SIFT/SU
Water



Classification algorithm

Object detection and/or Identification

YOLO



Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

Image Pre-processing

Image scaling
Contrast enhancement
Cropping and padding
Noise reduction




Feature extraction

Lines/Edge detection
Corners/Blobs/Points
(SIFT/SURF, PCA, Watershed)



Segmentation

Faster R-CNN



Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (2/2)

Image Pre-processing

Image scaling
Contrast enhancement
Cropping and padding
Noise reduction



Feature extraction

Lines/Edge detection
Corners/Blobs/Points
(SIFT/SURF, PCA, Watershed)



Classification algorithm

CNN



Object detection and/or Identification

YOLO




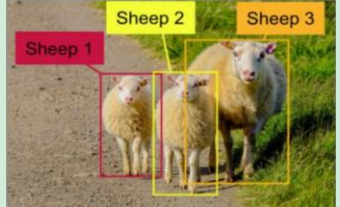
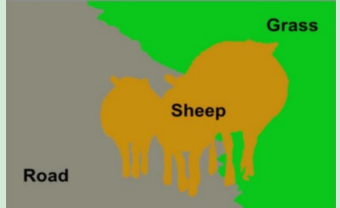
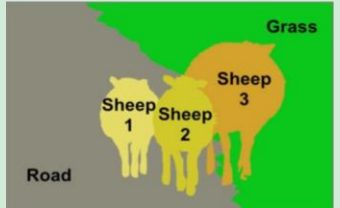
- We can start our process by applying image pre-processing (for instance if we have few data)
- Later we can use detection algorithm in order to find the object of interest
- Finally we can leverage feature extraction to post process the results to make them more precise or follow some field rules

AGENDA



1. Computer vision approaches
2. **An overview of the main Deep Learning algorithms**
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm
5. Resources to help you start

The 3 main approaches to machine learning vision are Classification, Detection and Segmentation

✂️ Approach	🎯 Objective	📁 Algorithms	👤 Examples
Classification	Associate a class, defining its content, to each image	<ul style="list-style-type: none"> CNN 	
Detection	Spot objects in an image according to predefined categories	<ul style="list-style-type: none"> YOLO SSD Faster R-CNN RFCN 	
Segmentation	Semantic: Link a category to each pixel in an image. Recognize a set of pixels that form distinct categories	<ul style="list-style-type: none"> Unet PSPNet ICNet DeepLab 	
	Per instance: Associate each pixel with the instance to which it belongs.	<ul style="list-style-type: none"> Mask R-CNN ResNet 	

These different approaches are based on **two different types of annotation**




Polygon

- Polygon annotations select all pixels belonging to an instance. It is used in segmentation



Bounding box

- Bounding box annotations give position, size and type of object contained in a rectangle parallel to the axes: used in detection






AGENDA



1. Computer vision approaches
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm
5. Resources to help you start

To train these models several **datasets are available online**

 Name	 Type	 Annotation
COCO	Natural landscapes	Instance segmentation
ADE20K		
SUN		Bounding boxes detection
IMAGE NET		
PASCAL VOC2012	City landscapes	Instance segmentation
CITYSCAPES		

- Is the **context** of the dataset close to context of my images?
- Does the **format** of the dataset compatible with my algorithm?
- What **classes** are labeled in the dataset?
- Is the dataset **balanced** between classes?

AGENDA



1. Computer vision approaches
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. **Measure the performance to improve your algorithm**
5. Resources to help you start

Several **metrics** can be used in computer vision depending on the objectives to reach

⚙️ Local Metrics


▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$

We set a threshold $\varepsilon \in [0,1]$

Prediction are considered true if $IoU \geq \varepsilon$.



Model Metrics

Several **metrics** can be used in computer vision depending on the objectives to reach

⚙️ Local Metrics


▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$

We set a threshold $\varepsilon \in [0,1]$

Prediction are considered true if $IoU \geq \varepsilon$.



Model Metrics

▪ Mean Intersect over Union (mIoU)

For a **fixed minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: over/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret

Several **metrics** can be used in computer vision depending on the objectives to reach

⚙️ Local Metrics


▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$

We set a threshold $\varepsilon \in [0,1]$

Prediction are considered true if $IoU \geq \varepsilon$.



▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**

$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**



Model Metrics

▪ Mean Intersect over Union (mIoU)

For a **fixed minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: over/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret


Several **metrics** can be used in computer vision depending on the objectives to reach

⚙️ Local Metrics

▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$



We set a threshold $\varepsilon \in [0,1]$

Prediction are considered true if $IoU \geq \varepsilon$.

▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**


$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**

▪ Precision-recall curve :

Precision and recall are calculated for different confidence levels (proba of the class model) in order to obtain the curve:

$$\text{precision} = f(\text{recall})$$



Model Metrics

▪ Mean Intersect over Union (mIoU)

For a **fixed minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: over/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret


Several **metrics** can be used in computer vision depending on the objectives to reach

⚙️ Local Metrics

▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$



We set a threshold $\varepsilon \in [0,1]$

Prediction are considered true if $IoU \geq \varepsilon$.

▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**


$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**

▪ Precision-recall curve :

Precision and recall are calculated for different confidence levels (proba of the class model) in order to obtain the curve:

$$\text{precision} = f(\text{recall})$$



Model Metrics

▪ Mean Intersect over Union (mIoU)

For a **fixed minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: over/under-prediction?
- Depends on a confidence threshold set

- Very "visual" metric, easy to interpret

▪ Average Precision (AP)

Once the IoU threshold is set, the AP of a class can be calculated as a measure of the area below the precision-recall curve.




- Difficult to interpret
- Depends on a fixed IoU threshold

- Independent of confidence level
- Gives an indication of the precision/recall trade-off

➤ **The mean Average Precision (mAP) is the AP averaged over all classes**

The AP reflects the ability to combine precision and recall for a model

Example n°1




$$AP \approx 0.74$$

Interpretation :

- We see that we can maintain a good trade-off between the level of precision and the level of recall
- The **green dot** guarantees an accuracy of about 80% for a 75% recall: most objects are spotted without too many errors

➤ An important AP ensures a certain efficiency of the model that will combine precision and recall

Example n°2



$$AP \approx 0.28$$

Interpretation :

We see here that we cannot guarantee precision and recall: we must choose to favour one or the other

- The **green zone** guarantees a precision > 75% accuracy but also a recall < 20% : we don't make much mistake but we spot very few objects
- The **pink zone** guarantees a recall > 80% but a precision < 5% : We spot almost all the object but most of the objects identified are not the one we want

➤ A low AP means it is difficult to reconcile accuracy and recall: one should be favored

AGENDA



1. Computer vision approaches
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm
5. **Resources to help you start**

Resources: all classical Machine Learning might come handy and a good understanding of Computer Vision libraries will be helpful

Computer Vision libraries



*Standard library for
deep learning*



*Standard library classical
computer vision*



*An alternative to
Pytorch*

ML + Viz libraries



*To develop a wide
range of ML models*



*To exploite model's
output and
aggregate them*



*To efficiently develop a
dashboard / front-end*

To use a pretrained detection model in Pytorch, several steps need to be done such as the download of the model, the analysis of the image by the model, and the printing of the prediction on the image



Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model

To use a **pretrained detection model** in Pytorch, several steps need to be done such as the **download** of the **model**, the **analysis** of the **image** by the model, and the **printing** of the **prediction** on the image



Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model

```
from PIL import Image
from torchvision import models
from torchvision import transforms as T
import matplotlib.pyplot as plt
import cv2
```

Example of code to import libraries

To use a **pretrained detection model** in Pytorch, several steps need to be done such as the **download** of the **model**, the **analysis** of the **image** by the model, and the **printing** of the **prediction** on the image



Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model

```
# Loading the model and the dataset
# Loads pretrained VGG model and sets it to eval mode

model = models.detection.fasterrcnn_resnet50_fpn(pretrained=True)
model = model.eval()
```

Example of code to download a detection model

To use a **pretrained detection model** in Pytorch, several steps need to be done such as the **download** of the **model**, the **analysis** of the **image** by the model, and the **printing** of the **prediction** on the image



Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model

```
COCO_INSTANCE_CATEGORY_NAMES = [  
    "__background__",  
    "person",  
    "bicycle",  
    "car",  
    "motorcycle",  
    "airplane",  
    "bus",  
    "train",  
    "truck",  
    "boat",
```

Example of code to define the classes of the detection model

To use a **pretrained detection model** in Pytorch, several steps need to be done such as the **download** of the **model**, the **analysis** of the **image** by the model, and the **printing** of the **prediction** on the image



Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model


```
img = Image.open(img_path) # Load the image
transform = T.Compose([T.ToTensor()]) # Defing PyTorch Transform
img = transform(img) # Apply the transform to the image
```


Example of code to process the input image


To use a pretrained detection model in Pytorch, several steps need to be done such as the download of the model, the analysis of the image by the model, and the printing of the prediction on the image


Illustration of the different steps


 Import libraries

 Download pretrained model

 Define the classes of the pretrained model

 Process the input image to the right format

 Predict the objects present in the image

 Plot the results of the prediction model


```
pred = model([img]) # Pass the image to the model
pred_class = [
    COCO_INSTANCE_CATEGORY_NAMES[i]
    for i in list(pred[0]["labels"].numpy())
] # Get the Prediction Score
pred_boxes = [
    [(i[0], i[1]), (i[2], i[3])]
    for i in list(pred[0]["boxes"].detach().numpy())
] # Bounding boxes
pred_score = list(pred[0]["scores"].detach().numpy())
pred_t = [pred_score.index(x) for x in pred_score if x > threshold][
    -1
] # Get list of index with score greater than threshold.
pred_boxes = pred_boxes[: pred_t + 1]
pred_class = pred_class[: pred_t + 1]
```


Example of code to predict the objects in the image


To use a **pretrained detection model** in Pytorch, several steps need to be done such as the **download** of the **model**, the **analysis** of the **image** by the model, and the **printing** of the **prediction** on the image


Illustration of the different steps


 Import libraries

 Download pretrained model

 Define the classes of the pretrained model

 Process the input image to the right format

 Predict the objects present in the image

 Plot the results of the prediction model

```
img = cv2.imread(img_path) # Read image with cv2
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) # Convert to RGB
for i in range(len(bboxes)):
    cv2.rectangle(
        img, bboxes[i][0], bboxes[i][1], color=(0, 255, 0), thickness=rect_th
    ) # Draw Rectangle with the coordinates
    cv2.putText(
        img,
        pred_cls[i],
        bboxes[i][0],
        cv2.FONT_HERSHEY_SIMPLEX,
        text_size,
        (0, 255, 0),
        thickness=text_th,
    ) # Write the prediction class
```

Example of code to plot the detected objects on the image

To use a pretrained detection model in Pytorch, several steps need to be done such as the download of the model, the analysis of the image by the model, and the printing of the prediction on the image

Illustration of the different steps



Import libraries



Download pretrained model



Define the classes of the pretrained model



Process the input image to the right format



Predict the objects present in the image



Plot the results of the prediction model



This tutorial is in the python file:
Tutorial_Computer_Vision_Detection.py

It can be used as a baseline for your project